Talent exists everywhere; opportunity does not. This belief has guided my commitment to diversity, equity, and inclusion throughout my academic career. As a Chinese international student and a woman in computer science, I have experienced the transformative power of access. I completed my undergraduate studies at a small liberal arts institution with three CS faculty and no research infrastructure, where I made the most of limited resources and took eight courses with the same professor to fulfill my major requirements. A summer internship at a major research institute revealed a world of possibilities I had never known existed. These experiences taught me that when we broaden access and challenge implicit biases, we unlock potential that might otherwise remain unseen. They also shaped my conviction that creating such opportunities for others is both a privilege and a responsibility.

At UW, I have worked to dismantle these barriers through advocacy and structural change. Advancing AI research requires not only technical excellence but also **intentional efforts to expand access, cultivate inclusive environments, and ensure our work serves all of humanity**. My commitment is reflected in measurable impact: mentoring 26 researchers (including 11 women, 23 people of color, and 18 first generation immigrants) across 10 institutions who have published at top venues; leading admissions, recruiting, and diversity efforts that shaped processes affecting hundreds of applicants; and developing research that centers pluralistic values, cultural inclusion, and global communities that are often overlooked in mainstream AI development.

1 Expanding Opportunity Through Admissions and Mentorship

Promoting Holistic Evaluation in Admissions For three consecutive years (2021–2023), I served as an Area Chair and Reviewer for Ph.D. admissions at UW CSE, managing the distribution and review of over 200 applications annually. Recognizing that admissions decisions profoundly shape our field's future, I actively promoted holistic evaluation beyond traditional metrics. I provided detailed feedback to reviewers when bias emerged in evaluations and deliberately surfaced strong candidates from non-traditional backgrounds in shortlist discussions, advocating for applicants whose potential might otherwise be overlooked by conventional measures.

Building Pathways for Underrepresented Students As a Mentor for the Pre-Application Mentorship Service (PAMS, 2021), I guided applicants from historically underrepresented backgrounds in preparing research statements and Ph.D. applications. I worked to expand the program's reach to international students through social and professional networks, recognizing the barriers many face in accessing graduate education. This work reinforced my commitment to creating multiple pathways into our field.

Mentoring Students from Diverse Backgrounds I have mentored 26 junior researchers across 10 institutions, with particular focus on supporting students who may not arrive with the strongest traditional credentials but demonstrate genuine enthusiasm and dedication. Having entered NLP/AI research only in my second year of Ph.D., I understand the challenges facing students from non-traditional paths. Several mentees have published their first papers and developed into long-term collaborators. For example, Kavel Rao entered research with limited background but through structured mentorship spanning three years, co-authored four papers at major venues and won the 2024 UW CSE Best Senior Thesis Award for our collaborative work on defeasible moral reasoning. Similarly, Kelly Chiu transitioned from a Master's program in Linguistics with a Psychology background into NLP research through our collaboration, resulting in three publications including two first-author papers at ICLR 2025 (Spotlight, Top 5.1%) and ACL 2025, and ultimately a Pre-Doctoral Researcher position at NYU. These examples demonstrate that with proper support systems, students from diverse backgrounds and non-traditional paths achieve excellence.

2 Fostering Inclusive Departmental Culture

Driving Institutional Change As the Student Representative on the UW Allen School Diversity and Inclusion Committee (2022), I reviewed diversity statements for the faculty search and helped ensure that candidates demonstrated a genuine commitment to fostering inclusive and equitable academic environments. As Co-Chair of the Prospective Student Committee (2022–2023), I redesigned visit-day programming to actively showcase departmental culture and provide admitted students from underrepresented backgrounds with clear pathways to community integration. I also co-organized the 2022 UW NLP Retreat, convening more than 200 participants to strengthen cross-lab networks and foster belonging across the research community. As Liaison for Faculty Recruiting (2022–2024), I ensured that diverse student perspectives were systematically incorporated into hiring decisions, recognizing that faculty composition profoundly shapes the culture and opportunities available to all students.

Creating Equitable Learning Environments As Head TA for CSE 447/517 with 230+ students, I worked to create an inclusive classroom through accessible materials, responsive support channels, and equitable evaluation. I prioritized clear communication and created structures, such as responsive Ed Discussion forums, that allowed students from diverse backgrounds and learning styles to succeed.

3 Research Serving Diverse Global Communities

My research advances equity by developing AI systems grounded in pluralistic human values and global cultural perspectives. Through work such as Kaleido (Oral at AAAI 2024) [1] and our position paper "A Roadmap to Pluralistic Alignment" (ICML 2024) [2], now recognized as foundational in this area, I champion the view that multiple and sometimes conflicting values can all be legitimate, rather than relying on uniform approaches that may marginalize minority perspectives. This commitment extends to my work on CulturalBench [3] and Culture-Gen [4], which address the systematic underrepresentation of non Western cultural knowledge in AI systems, and to multilingual safety efforts like PolyGuard [5], which expands AI safety tools to seventeen languages. Through my position paper on political neutrality in AI (Oral at ICML 2025) [6], my colleagues and I also develop frameworks for building systems that support democratic discourse by avoiding the imposition of dominant viewpoints.

4 Broadening Participation Through Public Engagement

Inspiring the Next Generation I am committed to making AI accessible to young learners and inspiring the next generation of diverse researchers. In 2022 and 2023, I delivered guest lectures on "Can We Teach Machines Human Ethics and Values?" to high school students at The Downtown School in Seattle, introducing complex AI concepts in ways that were engaging and age-appropriate. I have also mentored individual high school students, including Abhay Gupta (John Jay Senior High School), who reached out seeking early research experience in AI.

Communicating with Broad Audiences I have actively engaged with media to make AI research accessible to the public, recognizing that informed discourse requires clear communication across communities. My work has appeared in outlets including The New York Times, The New Yorker, Vox, Nature Outlook, Wired, IEEE Spectrum, and The Guardian, where I provide accessible explanations that help diverse audiences understand AI's capabilities, limitations, and societal implications.

Future Commitments

Institutional Leadership I plan to serve actively on graduate admissions and faculty hiring committees, promoting holistic evaluation and helping recruit scholars whose excellence may not be reflected in traditional metrics. I am committed to fostering a fair, student-empowering culture by strengthening mentoring structures, advancing equitable evaluation practices, and helping shape curricula that prepare students from all backgrounds for meaningful engagement with AI.

Building an Inclusive Research Group I will create a lab culture grounded in respect, psychological safety, and belonging where students from diverse backgrounds can thrive. This includes structured onboarding, peer mentoring networks, clear communication of expectations, and proactive outreach to underrepresented communities. I will ensure my mentoring addresses each student's unique needs and aspirations while maintaining high standards for all.

Expanding Mentorship at Scale I will create mentorship programs extending beyond my research group, including programs for motivated students from underrepresented backgrounds and publicly accessible resources to broaden participation. I will share teaching materials, recorded lectures, and research guidance openly to reduce barriers for students without access to elite institutions.

Research with Global Impact I will continue developing AI systems that serve pluralistic values and diverse global communities rather than privileged subpopulations. This includes research on cross-cultural AI alignment, multilingual safety, and frameworks that preserve rather than homogenize human diversity [7].

Through these efforts, I aim to build research communities where excellence is recognized in its many forms, where students from all backgrounds can contribute and thrive, and where our technical innovations genuinely serve humanity's diversity. **Diversity is not ancillary to my academic mission; it is indispensable to building AI systems that honor humanity's plurality and advance equitable outcomes.**

* and † denote equal contribution.

References

[1] Taylor Sorensen, **Liwei Jiang**, Jena Hwang, Sydney Levine, Valentina Pyatkin, Peter West, Nouha Dziri, Ximing Lu, Kavel Rao, Chandra Bhagavatula, Maarten Sap, John Tasioulas, and Yejin Choi. Value Kaleidoscope: Engaging AI with Pluralistic Human Values, Rights, and Duties.
AAAI 2024 (Oral, Top 3%).

[2] Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, **Liwei Jiang**, Ximing Lu, Nouha Dziri, Tim Althoff, and Yejin Choi. Position Paper: A Roadmap to Pluralistic Alignment.
ICML 2024.

[3] Yu Ying Chiu, **Liwei Jiang**, Bill Yuchen Lin, Chan Young Park, Shuyue Stella Li, Sahithya Ravi, Mehar Bhatia, Maria Antoniak, Yulia Tsvetkov, Vered Shwartz, and Yejin Choi. CulturalBench: A Robust, Diverse, and Challenging Cultural Benchmark by Human-AI CulturalTeaming.
ACL 2025.

[4] Huihan Li, **Liwei Jiang**, Jena D. Huang, Hyunwoo Kim, Sebastin Santy, Taylor Sorensen, Bill Yuchen Lin, Nouha Dziri, Xiang Ren, and Yejin Choi. Culture-Gen: Revealing Global Cultural Perception in Language Models through Natural Language Prompting.
COLM 2024.

[5] Priyanshu Kumar*, Devansh Jain*, Akhila Yerukola, **Liwei Jiang**, Himanshu Beniwal, Tom Hartvigsen, and Maarten Sap. PolyGuard: A Multilingual Safety Moderation Tool for 17 Languages.
COLM 2025.

[6] Jillian Fisher, Ruth Elisabeth Appel, Chan Young Park, Yujin Potter, **Liwei Jiang**, Taylor Sorensen, Shangbin Feng, Yulia Tsvetkov, Margaret Roberts, Jennifer Pan, Dawn Song, and Yejin Choi. Position Paper: Political Neutrality in AI is Impossible–But Here's How to Approximate It.
ICML 2025 (Oral, Top 3.3%).

[7] **Liwei Jiang**, Yuanjun Chai, Margaret Li, Mickel Liu, Raymond Fok, Nouha Dziri, Yulia Tsvetkov, Maarten Sap, Alon Albalak, and Yejin Choi. Artificial Hiveminds: The Open-Ended Homogeneity of Language Models (and Beyond).
NeurIPS 2025, Datasets & Benchmarks Track (Best Paper Award, Oral).